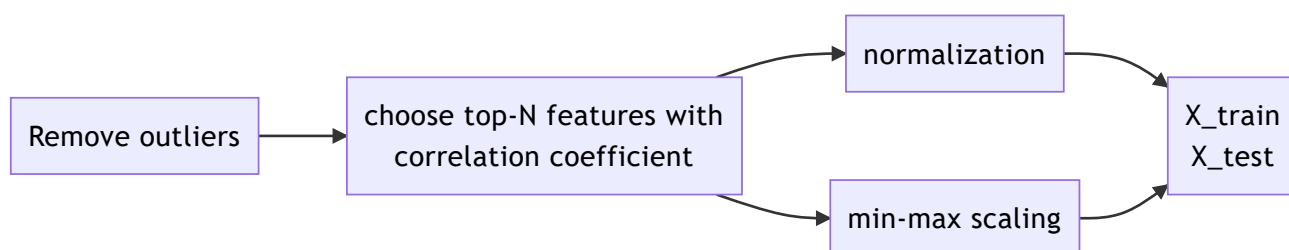


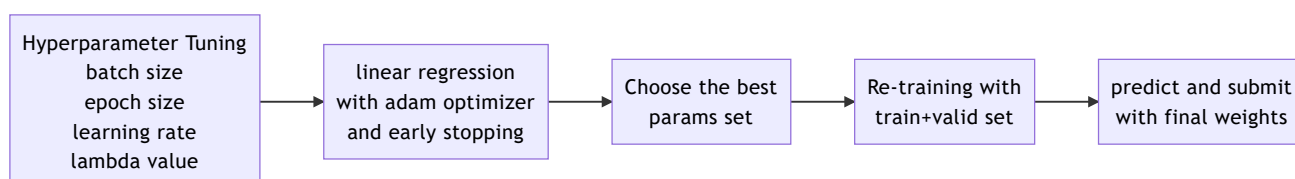
# EEML HW1 - Programming Report

B06702064 會計五 林聖硯

## Data preprocessing pipeline



## Modeling pipeline



## Tasks' details

任務	說明
Remove outliers	使用outliers的定義 $[\mu - 1.5IQR, \mu + 1.5IQR]$ 移除y值(第十小時PM2.5)過高之observation
choose top-N features with correlation coefficient (N = <b>15</b> by try and error)	使用Pearson correlation coefficient 選出與y最高的前N個features
Hyperparameter tuning	Grid search 以下參數組合 batch size= [64, 128, 256, 512, 1024] epoch size = [20, 30, 50, 100] learning rate = [0.002, 0.005, 0.01, 0.025, 0.05] lambda= [0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1, 5]

請實作以下兩種不同feature的模型，回答第 1 ~ 2 題：

- **Dataset1**: 抽全部9小時內的污染源feature當作一次項(加bias)
- **Dataset2**: 抽全部9小時內pm2.5的一次項當作feature(加bias)
- 以上兩個dataset，除了**choose top-N features with correlation coefficient**此流程外，都有經過**其他所有的process**(包含調整參數等等)，以下問題將由模型調參後的結果回答


1.記錄誤差值 (RMSE)(根據kaggle public+private分數)，討論兩種feature的影響

dataset	Tuning result	validation RMSE	public RMSE	private RMSE
Dataset1	Batch Size=128 Epoch Size=100 (running 93 epoch) learning rate=0.005 lambda=0.001	0.655284	16.20125	17.11370
Dataset2	Batch Size=256 Epoch Size=100 (running 52 epoch) learning rate=0.025 lambda=0.001	4.735835	5.23247	4.87396

### Observation:

從第一個Dataset上可發現，使用全部的features當做資料進行訓練，雖然在validation set上表現得很好，但卻沒有辦法generalize到public test/private test set上，顯示模型因為過多的feature而有**overfitting**的情況出現。而在Dataset2上，雖然validation RMSE的結果沒有那麼好，但模型generalize的能力相較Dataset1就強上許多。可能的原因是PM2.5本來就是一個有週期性的時間序列資料，所以只要模型有學出這件事情，用前九個小時的PM2.5來預估第十小時的PM2.5自然準確度就會高，加上其他features對模型來說就是增加額外的noise，反而會導致模型學不好這件事。

移除outlier後，PM2.5在不同小時畫出來之折線圖



2. 解釋什麼樣的data preprocessing可以improve你的training/testing accuracy，e.g., 你怎麼挑掉你覺得不適合的data points。請提供數據(RMSE)以佐證你的想法。

- 首先，我對y值進行EDA，並且發現了在PM2.5=1000時以及PM2.5=100時有兩個斷層。故以下的資料處理步驟，我有針對這兩點特別進行處理。

```
print(data_train[data_train['PM2.5'] > 5000].shape)
• print(data_train[data_train['PM2.5'] > 4000].shape)
print(data_train[data_train['PM2.5'] > 3000].shape)
print(data_train[data_train['PM2.5'] > 2000].shape)
print(data_train[data_train['PM2.5'] > 1000].shape)
#斷層出現在這裡
print(data_train[data_train['PM2.5'] > 800].shape)
print(data_train[data_train['PM2.5'] > 600].shape)
print(data_train[data_train['PM2.5'] > 400].shape)
print(data_train[data_train['PM2.5'] > 200].shape)
print(data_train[data_train['PM2.5'] > 100].shape)
#第二個斷層
print(data_train[data_train['PM2.5'] > 50].shape)
```

✓ 0.5s

```
(3, 18)
(3, 18)
(3, 18)
(5, 18)
(5, 18)
(32, 18)
(86, 18)
(86, 18)
(87, 18)
(98, 18)
(967, 18)
```

- 接著我分別對模型進行以下的資料前處理之組合

資料處理之不同組合	validation RMSE	public RMSE	private RMSE
不進行任何處理	0.655284	16.20125	17.11370
移除y>1000的observation	1.322826	16.06233	16.80802
移除y為outlier之observation	1.422711	16.30997	17.18779
移除y為outlier之observation + normalization	0.655284	16.20125	17.11370
移除y為outlier之observation + min-max scaling	0.683863	16.12268	16.98989
移除y為outlier之observation + choose top 15 features with correlation coefficient	4.94098	5.42117	4.91746
移除y為outlier之observation + choose top 15 features with correlation coefficient + normalization	4.581977	5.48004	4.78525
移除y為outlier之observation + choose top 15 features with correlation coefficient + min-max scaling	4.657251	4.97667	4.76087

## Observation

- 從上述圖表中可以看出即使移除outliers加上normalization或min-max scaling，都會讓整個模型overfit validation dataset
- 真正能讓模型在out-sample中大幅進步的關鍵因素為**去掉不重要的features**，做了這步之後，我的模型基本上在training、validation、public/private的資料集上表現不會差太多。原因如上所述，有一些不重要的features加進來訓練可能會增加整個任務的noise，導致模型沒辦法正確的學習

## 3.一些模型訓練的心得

其實我是public dataset的第一名，我最終選了兩個public score最高的submission繳交成private，結果在private名次竟然掉到了十一名。其實public score最高的兩個model都是我在意外之中(模型還沒完全建立的過程中)做出來的，我還很懊悔沒辦法reproduce他們。結果我明明

有private score更好的模型(也是在模型架好之後做的)，只是因為他們的public score不夠高，我就沒有選擇他們。真的不能完全依賴public score來判斷模型好壞，也要考慮你在model training的時候多了或少了什麼步驟。

Submission and Description	Private Score	Public Score	Use for Final Score
<a href="#">submission_1013_2.csv</a> 9 days ago by <a href="#">b06702064_Martin</a> 1013_2	5.12013	4.64726	<input checked="" type="checkbox"/>
<a href="#">submission_1014_4.csv</a> 8 days ago by <a href="#">b06702064_Martin</a> 1014 remove outliers, top 15 corr features / tuning batch size, epoch size, learning rate / L2 normalization	4.84971	4.84708	<input checked="" type="checkbox"/>
<a href="#">submission_1020_3.csv</a> 2 days ago by <a href="#">b06702064_Martin</a> 1020 remove outliers, top 15 corr features, normalization / tuning batch size, epoch size, learning rate / randomize weight and bias, early stopping, do not train with all data finally / L2 normalization	4.70913	4.89540	<input type="checkbox"/>
<a href="#">submission_1020_5.csv</a> 2 days ago by <a href="#">b06702064_Martin</a> 1020 remove outliers, top 15 corr features, min-max scaling / tuning batch size, epoch size, learning rate / randomize weight and bias, early stopping, do not train with all data finally / L2 normalization	4.73896	4.91783	<input type="checkbox"/>
<a href="#">submission_1016_1.csv</a> 6 days ago by <a href="#">b06702064_Martin</a> 1016 remove outliers, top 15 corr features / tuning batch size, epoch size, learning rate / do not traing with all data / L2 normalization	4.70995	4.95081	<input type="checkbox"/>

## Problem 1 - Logistic Regression

(a)

$$\begin{aligned} P_{w,b}(C_1|x) &= \sigma\left(\sum_{i=1}^4 w_i x_i + b\right) \\ &= \sigma(-7 + 0 - 3 + 50 + 1) \\ &= \frac{1}{1 + \exp(-41)} \\ &\approx 1 \end{aligned}$$

(b)

$$\begin{aligned} w^*, b^* &= \arg \max_{w,b} L(w, b) \iff w^*, b^* = \arg \min_{w,b} -L(w, b) \\ -\ln L(w, b) &= -\ln f_{w,b}(x^1) - \ln f_{w,b}(x^2) - \ln(1 - f_{w,b}(x^3)) \cdots - \ln f_{w,b}(x^N) \\ &= -[1 \times \ln f_{w,b}(x^1) + 0 \times \ln(1 - f_{w,b}(x^1))] - [1 \times \ln f_{w,b}(x^2) + 0 \times \ln(1 - f_{w,b}(x^2))] \\ &\quad \dots [1 \times \ln f_{w,b}(x^N) + 0 \times \ln(1 - f_{w,b}(x^N))] \\ &= \sum_{n=1}^N -[\hat{y}^n \ln f_{w,b}(x^n) + (1 - \hat{y}^n) \ln(1 - f_{w,b}(x^n))] \end{aligned}$$

The above equation can be viewed as calculating cross entropy between two Bernoulli distribution  $P$  and  $Q$ , where  $\text{Prob}(x^n \in \text{class1}|P) = \hat{y}^n$ ,  $\text{Prob}(x^n \in \text{class2}|P) = 1 - \hat{y}^n$  and  $\text{Prob}(x^n \in \text{class1}|Q) = f(x^n)$ ,  $\text{Prob}(x^n \in \text{class2}|Q) = 1 - f(x^n)$

(c)

$$\begin{aligned} \frac{\partial \ln f_{w,b}(x)}{\partial w_i} &= \frac{\partial \ln f_{w,b}(x)}{\partial z} \times \frac{\partial z}{\partial w_i}, \text{ where } f_{w,b}(x) = \sigma(z), z = \sum_{i=1}^4 x_i w_i + b \\ &= \frac{1}{\sigma(z)} \times \frac{\partial \sigma(z)}{\partial z} \times x_i \\ &= \frac{1}{\sigma(z)} \times \sigma(z)[1 - \sigma(z)] \times x_i \\ &= [1 - \sigma(z)] x_i \\ \frac{\partial \ln(1 - f_{w,b}(x))}{\partial w_i} &= \frac{\partial \ln(1 - \sigma(z))}{\partial z} \times \frac{\partial z}{\partial w_i} \\ &= -\frac{1}{1 - \sigma(z)} \times \sigma(z)[1 - \sigma(z)] \times x_i \\ &= -\sigma(z) x_i \end{aligned}$$

$$\begin{aligned}
 \Rightarrow -\frac{\partial \ln L(w, b)}{\partial w_i} &= \sum_{n=1}^N -[\hat{y}^n \frac{\partial \ln f_{w,b}(x^n)}{\partial w_i} + (1 - \hat{y}^n) \frac{\partial \ln (1 - f_{w,b}(x^n))}{\partial w_i}] \\
 &= \sum_{n=1}^N -[\hat{y}^n [1 - f_{w,b}(x^n)] x_i^n - (1 - \hat{y}^n) f_{w,b}(x^n) x_i^n] \\
 &= \sum_{n=1}^N -[\hat{y}^n - f_{w,b}(x^n)] x_i^n
 \end{aligned}$$

Thus, the update rule is  $w_i^{t+1} \leftarrow w_i^t - \eta_t \sum_{n=1}^N -[\hat{y}^n - f_{w,b}(x^n)] x_i^n$

## Problem 2 - Closed-Form Linear Regression Solution

(a) Since  $w \in \mathbb{R}^1$  in this example, denote  $w^T = w$

$$\begin{aligned}
 L_{ssq}(w, b) &= \frac{1}{2 \times 5} = \sum_{i=1}^5 (y_i - wx_i - b)^2 \\
 \frac{\partial L_{ssq}(w, b)}{\partial w} &= \frac{1}{5} \sum_{i=1}^5 (y_i - wx_i - b)(-x_i) = 0 \\
 \Rightarrow w \sum_{i=1}^5 x_i^2 + b \sum_{i=1}^5 x_i &= \sum_{i=1}^5 x_i y_i \\
 \frac{\partial L_{ssq}(w, b)}{\partial b} &= \frac{1}{5} \sum_{i=1}^5 (y_i - wx_i - b)(-1) = 0 \\
 \Rightarrow w \sum_{i=1}^5 x_i + 5b &= \sum_{i=1}^5 y_i
 \end{aligned}$$

From the two equations above, we can get

$$\begin{cases} 55w + 15b &= 59.7 \\ 15w + 5b &= 16.8 \end{cases}$$

$$\Rightarrow w = 0.93, b = 0.57$$

$$(b) \text{ Let } \beta = \begin{bmatrix} b \\ w_1 \\ w_2 \\ \vdots \\ w_k \end{bmatrix} \in \mathbb{R}^{k+1}, \mathbf{X} = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_N^T \end{bmatrix} \in \mathbb{R}^{N \times (k+1)}, \text{ where } x_n = \begin{bmatrix} 1 \\ x_{n,1} \\ x_{n,2} \\ \vdots \\ x_{n,k} \end{bmatrix} \in \mathbb{R}^{k+1}$$

$$\text{and } \mathbf{Y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \in \mathbb{R}^N$$

The linear regression model can be represented as  $\mathbf{X}\beta = \mathbf{Y}$

$$\begin{aligned} L(\beta) &= \frac{1}{2N}(\mathbf{X}\beta - \mathbf{Y})^T(\mathbf{X}\beta - \mathbf{Y}) \\ &= \frac{1}{2N}((\mathbf{X}\beta)^T - \mathbf{Y}^T)(\mathbf{X}\beta - \mathbf{Y}) \\ &= \frac{1}{2N}\beta^T \mathbf{X}^T \mathbf{X} \beta - 2(\mathbf{X}\beta)^T \mathbf{Y} - \mathbf{Y}^T \mathbf{Y} \\ \frac{\partial L(\beta)}{\partial \beta} &= \frac{1}{2N}[\mathbf{X}^T \mathbf{X} \frac{\partial \beta^T \beta}{\partial \beta} - 2\mathbf{X}^T \mathbf{Y} \frac{\partial \beta^T}{\partial \beta}] = 0 \\ &\implies (\mathbf{X}^T \mathbf{X})\beta = \mathbf{X}^T \mathbf{Y} \end{aligned}$$

Assume  $\mathbf{X}^T \mathbf{X}$  is invertible,

$$\implies \beta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

(c) Following the notation of 2(b)

$$\begin{aligned} L_{reg}(w, b) &= \frac{1}{2N} \sum_{i=1}^N (y_i - (\mathbf{w}^T \mathbf{x}_i + b))^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2 \\ &= \frac{1}{2N} (\mathbf{X}\beta - \mathbf{Y})^T (\mathbf{X}\beta - \mathbf{Y}) + \frac{\lambda}{2} \beta_{-1}^T \beta_{-1}, \text{ where } \beta_{-1} = \begin{bmatrix} 0 \\ w_1 \\ w_2 \\ \vdots \\ w_k \end{bmatrix} \in \mathbb{R}^{k+1} \end{aligned}$$



Compute  $\frac{\partial \beta_{-1}^T \beta_{-1}}{\partial \beta}$

$$\begin{aligned} \frac{\partial \beta_{-1}^T \beta_{-1}}{\partial \beta} &= \frac{\partial \beta_{-1}}{\partial \beta} \beta_{-1} + \frac{\partial \beta_{-1}}{\partial \beta} \beta_{-1} \\ &= 2 \times \begin{bmatrix} 0 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{bmatrix} \beta_{-1} \\ &\stackrel{\text{def}}{=} 2\mathbf{E}^* \beta_{-1}, \text{ where } \mathbf{E}^* \in \mathbb{R}^{(k+1) \times (k+1)} \end{aligned}$$

$$\begin{aligned} \frac{\partial L_{reg}(w, b)}{\partial \beta} &= \frac{1}{2N} [2\mathbf{X}^T \mathbf{X} \beta - 2\mathbf{X}^T \mathbf{Y}] + \lambda \mathbf{E}^* \beta = 0 \\ \implies \beta \left( \frac{1}{N} \mathbf{X}^T \mathbf{X} + \lambda \mathbf{E}^* \right) &= \frac{1}{N} \mathbf{X}^T \mathbf{Y} \end{aligned}$$

Assume  $\frac{1}{N} \mathbf{X}^T \mathbf{X} + \lambda \mathbf{E}^*$  is invertible,

$$\implies \beta = \left( \frac{1}{N} \mathbf{X}^T \mathbf{X} + \lambda \mathbf{E}^* \right)^{-1} \frac{1}{N} \mathbf{X}^T \mathbf{Y}$$

### Problem3 - Noise and regularization

$$\begin{aligned} \tilde{L}_{ssq}(w, b) &= \mathbb{E} \left[ \frac{1}{2N} \sum_{i=1}^N [(\mathbf{w}^T (\mathbf{x}_i + \eta_i)) + b - y_i]^2 \right] \\ &= \frac{1}{2N} \mathbb{E} \left[ \sum_{i=1}^N [(\mathbf{w}^T \mathbf{x}_i + b - y_i) + \mathbf{w}^T \eta_i]^2 \right] \\ &= \frac{1}{2N} \sum_{i=1}^N \left( \mathbb{E} [(\mathbf{w}^T \mathbf{x}_i + b - y_i)^2] + 2(\mathbf{w}^T \mathbf{x}_i + b - y_i) \mathbb{E} [\mathbf{w}^T \eta_i] + \mathbb{E} [(\mathbf{w}^T \eta_i)^2] \right) \end{aligned}$$

Compute  $\mathbb{E} [\mathbf{w}^T \eta_i]$

$$\mathbb{E} [\mathbf{w}^T \eta_i] = \sum_{j=1}^k w_j \mathbb{E} [\eta_{i,j}] = 0$$

Compute  $\mathbb{E}[(\mathbf{w}^T \eta_i)^2]$

$$\begin{aligned}\mathbb{E}[(\mathbf{w}^T \eta_i)^2] &= \text{Var}[\mathbf{w}^T \eta_i] - (\mathbb{E}[\mathbf{w}^T \eta_i])^2 \\ &= w_1^2 \text{Var}[\eta_{i,1}] + w_2^2 \text{Var}[\eta_{i,2}] + \cdots + w_k^2 \text{Var}[\eta_{i,k}] \\ &= \sum_{j=1}^k w_j^2 \text{Var}[\eta_{i,j}] \\ &= \sum_{j=1}^k w_j^2 \sigma^2 \\ &= \sigma^2 \|\mathbf{w}\|^2\end{aligned}$$

$$\begin{aligned}\tilde{L}_{ssq}(w, b) &= \frac{1}{2N} \sum_{i=1}^N \left( \mathbb{E}[(\mathbf{w}^T \mathbf{x}_i + b - y_i)^2] + 2(\mathbf{w}^T \mathbf{x}_i + b - y_i) \mathbb{E}[\mathbf{w}^T \eta_i] + \mathbb{E}[(\mathbf{w}^T \eta_i)^2] \right) \\ &= \frac{1}{2N} \sum_{i=1}^N \left[ (f_{w,b}(x_i) - y_i)^2 + \sigma^2 \|\mathbf{w}\|^2 \right] \\ &= \frac{1}{2N} \left[ \sum_{i=1}^N (f_{w,b}(x_i) - y_i)^2 \right] + \frac{\sigma^2}{2} \|\mathbf{w}\|^2\end{aligned}$$

Thus, minimizing the expected sum-of-squares loss in the presence of input noise is equivalent to minimizing noise-free sum-of-squares loss with the addition of a  $L^2$ -regularization term on the weights.

## Problem4 - Kaggle Hacker

(a)

$$\begin{aligned} e_k &= \frac{1}{N} \sum_{i=1}^N (g_k(x_i) - y_i)^2 \\ &= \frac{1}{N} \left[ \sum_{i=1}^N (g_k(x_i))^2 - 2 \sum_{i=1}^N g_k(x_i) y_i + \sum_{i=1}^N y_i^2 \right] \\ &= s_k - \frac{2}{N} \sum_{i=1}^N g_k(x_i) y_i + e_0 \\ &\Rightarrow \frac{2}{N} \sum_{i=1}^N g_k(x_i) y_i = s_k + e_0 - e_k \\ &\Rightarrow \sum_{i=1}^N g_k(x_i) y_i = \frac{N}{2} (s_k + e_0 - e_k) \end{aligned}$$

(b) Let  $\mathbf{G} = \begin{bmatrix} g_1(x_1) & g_2(x_1) & \cdots & g_K(x_1) \\ g_1(x_2) & g_2(x_2) & \cdots & g_K(x_2) \\ \vdots & \vdots & \ddots & \vdots \\ g_1(x_N) & g_2(x_N) & \cdots & g_K(x_N) \end{bmatrix} \in \mathbb{R}^{N \times K}$ ,  $\alpha = \begin{bmatrix} \alpha_1 \\ \alpha_2 \\ \vdots \\ \alpha_K \end{bmatrix} \in \mathbb{R}^K$ ,  $\mathbf{Y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \in \mathbb{R}^N$

$$\begin{aligned} \min_{\alpha_1, \dots, \alpha_K} L_{test} \left( \sum_{k=1}^K \alpha_k g_k \right) &= \min_{\alpha_1, \dots, \alpha_K} \frac{1}{N} (\mathbf{G}\alpha - \mathbf{Y})^T (\mathbf{G}\alpha - \mathbf{Y}) \\ \frac{\partial L_{test}}{\partial \alpha} &= \frac{1}{N} [\mathbf{G}^T \mathbf{G} 2\alpha - 2\mathbf{G}^T \mathbf{Y}] = 0 \end{aligned}$$

Assume  $\mathbf{G}^T \mathbf{G}$  is invertible,

$$\Rightarrow \alpha = (\mathbf{G}^T \mathbf{G}^{-1}) \mathbf{G}^T \mathbf{Y}$$